

# CAP 4611 Programming Assignment 5: Neural Network Report

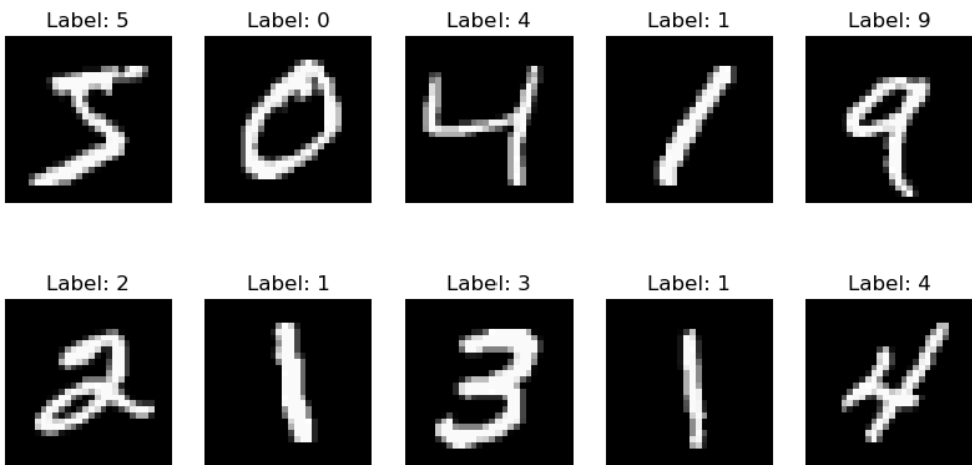
By: Kashyap Bathina (ka964126) Section 1

---

## Part 2: Building and Training a Neural Network Model

Results, graphs, reports:

(60000, 28, 28)  
Training data shape: (60000, 784)  
Test data shape: (10000, 784)



Model: "sequential\_2"

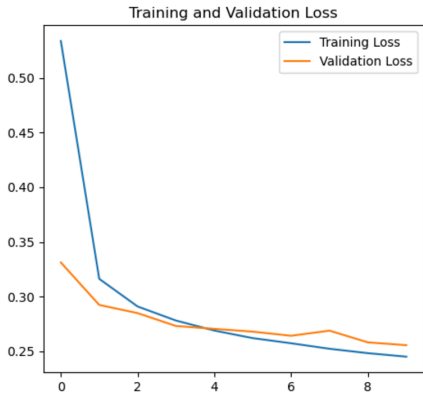
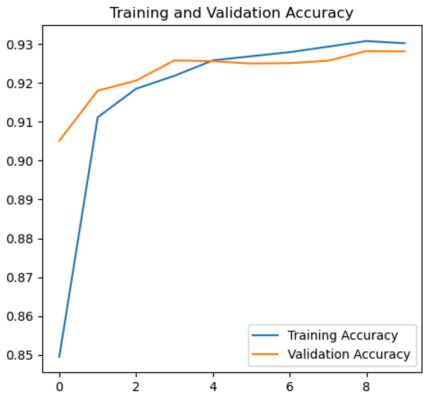
Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 8)	6,280
dense_5 (Dense)	(None, 10)	90

Total params: 6,370 (24.88 KB)

Trainable params: 6,370 (24.88 KB)

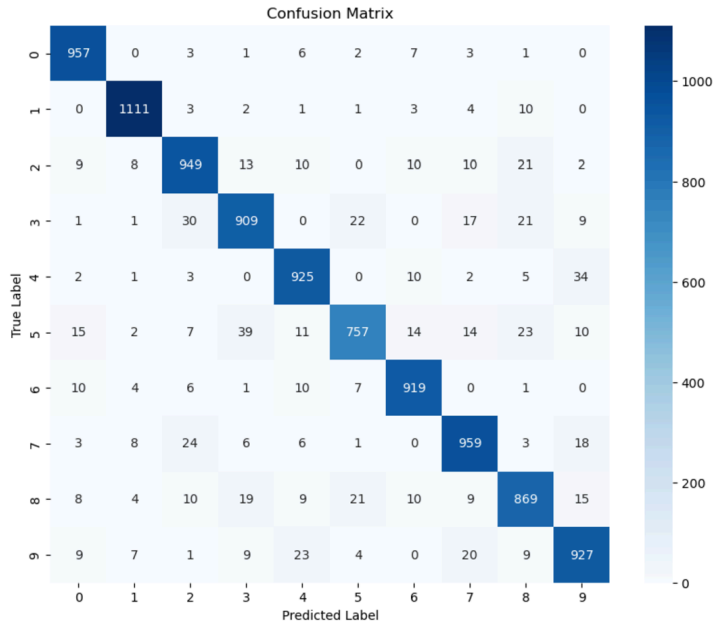
Non-trainable params: 0 (0.00 B)

Epoch 1/10  
1875/1875 2s 889us/step - accuracy: 0.7443 - loss: 0.8646 - val\_accuracy: 0.9051 - val\_loss: 0.3312  
Epoch 2/10  
1875/1875 2s 884us/step - accuracy: 0.9088 - loss: 0.3265 - val\_accuracy: 0.9180 - val\_loss: 0.2924  
Epoch 3/10  
1875/1875 2s 837us/step - accuracy: 0.9170 - loss: 0.2972 - val\_accuracy: 0.9206 - val\_loss: 0.2849  
Epoch 4/10  
1875/1875 2s 886us/step - accuracy: 0.9229 - loss: 0.2766 - val\_accuracy: 0.9258 - val\_loss: 0.2731  
Epoch 5/10  
1875/1875 2s 797us/step - accuracy: 0.9256 - loss: 0.2694 - val\_accuracy: 0.9256 - val\_loss: 0.2705  
Epoch 6/10  
1875/1875 2s 810us/step - accuracy: 0.9280 - loss: 0.2591 - val\_accuracy: 0.9250 - val\_loss: 0.2679  
Epoch 7/10  
1875/1875 2s 865us/step - accuracy: 0.9287 - loss: 0.2544 - val\_accuracy: 0.9251 - val\_loss: 0.2642  
Epoch 8/10  
1875/1875 2s 858us/step - accuracy: 0.9305 - loss: 0.2517 - val\_accuracy: 0.9257 - val\_loss: 0.2688  
Epoch 9/10  
1875/1875 2s 798us/step - accuracy: 0.9319 - loss: 0.2421 - val\_accuracy: 0.9282 - val\_loss: 0.2581  
Epoch 10/10  
1875/1875 2s 831us/step - accuracy: 0.9301 - loss: 0.2446 - val\_accuracy: 0.9281 - val\_loss: 0.2555



313/313 0s 856us/step - accuracy: 0.9190 - loss: 0.2929  
Test Accuracy: 0.9282  
313/313 0s 707us/step

	precision	recall	f1-score	support
0	0.94	0.98	0.96	980
1	0.97	0.98	0.97	1135
2	0.92	0.92	0.92	1032
3	0.91	0.90	0.90	1010
4	0.92	0.94	0.93	982
5	0.93	0.85	0.89	892
6	0.94	0.96	0.95	958
7	0.92	0.93	0.93	1028
8	0.90	0.89	0.90	974
9	0.91	0.92	0.92	1009
accuracy			0.93	10000
macro avg	0.93	0.93	0.93	10000
weighted avg	0.93	0.93	0.93	10000



## Part 3: Experimenting with a Larger Hidden Layer

Results, graphs, reports:

Model: "sequential\_5"

Layer (type)	Output Shape	Param #
dense_10 (Dense)	(None, 128)	100,480
dense_11 (Dense)	(None, 10)	1,290

Total params: 101,770 (397.54 KB)

Trainable params: 101,770 (397.54 KB)

Non-trainable params: 0 (0.00 B)

Epoch 1/10

1875/1875 — 3s 1ms/step - accuracy: 0.8786 - loss: 0.4312 - val\_accuracy: 0.9569 - val\_loss: 0.1443

Epoch 2/10

1875/1875 — 2s 1ms/step - accuracy: 0.9646 - loss: 0.1263 - val\_accuracy: 0.9680 - val\_loss: 0.1003

Epoch 3/10

1875/1875 — 2s 1ms/step - accuracy: 0.9772 - loss: 0.0804 - val\_accuracy: 0.9723 - val\_loss: 0.0853

Epoch 4/10

1875/1875 — 2s 1ms/step - accuracy: 0.9809 - loss: 0.0611 - val\_accuracy: 0.9753 - val\_loss: 0.0799

Epoch 5/10

1875/1875 — 3s 1ms/step - accuracy: 0.9864 - loss: 0.0462 - val\_accuracy: 0.9771 - val\_loss: 0.0701

Epoch 6/10

1875/1875 — 3s 1ms/step - accuracy: 0.9898 - loss: 0.0341 - val\_accuracy: 0.9736 - val\_loss: 0.0835

Epoch 7/10

1875/1875 — 3s 1ms/step - accuracy: 0.9915 - loss: 0.0275 - val\_accuracy: 0.9770 - val\_loss: 0.0794

Epoch 8/10

1875/1875 — 2s 1ms/step - accuracy: 0.9942 - loss: 0.0220 - val\_accuracy: 0.9782 - val\_loss: 0.0724

Epoch 9/10

1875/1875 — 2s 1ms/step - accuracy: 0.9951 - loss: 0.0173 - val\_accuracy: 0.9770 - val\_loss: 0.0839

Epoch 10/10

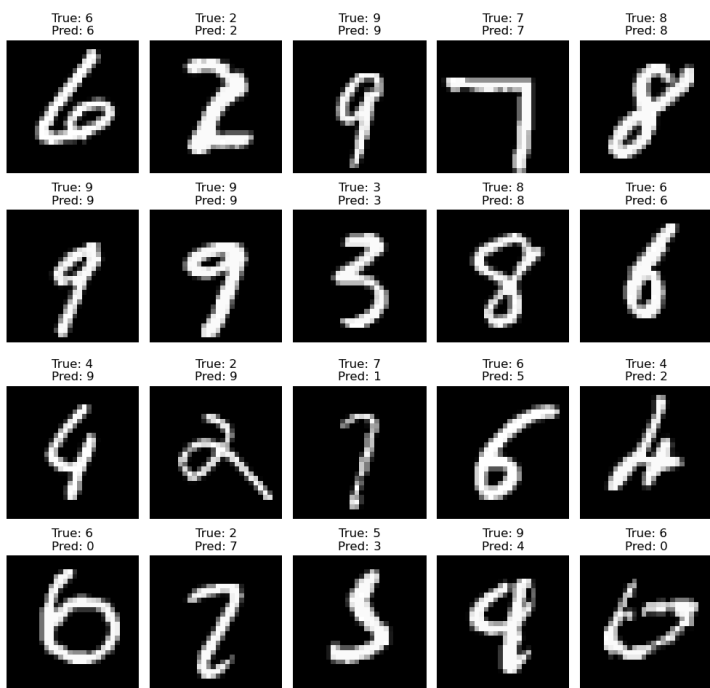
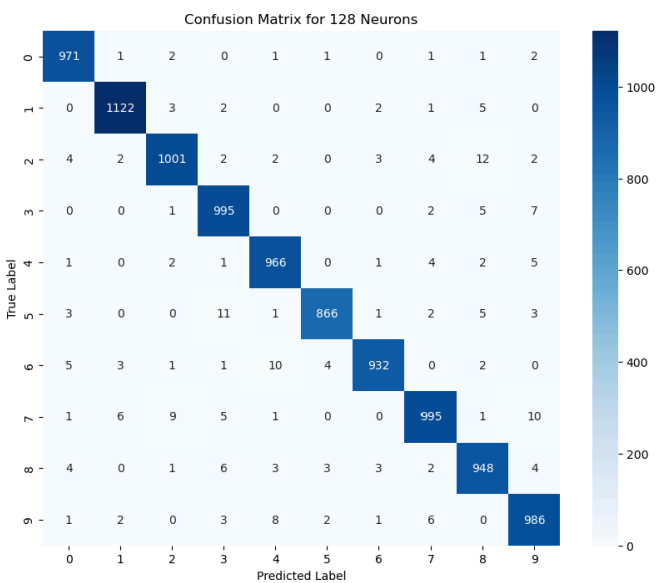
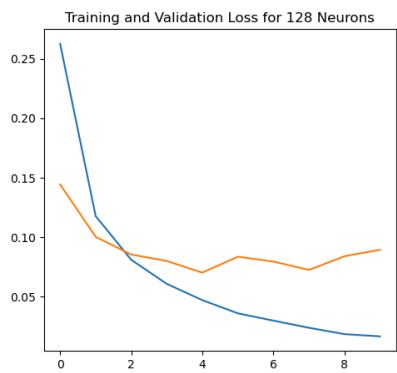
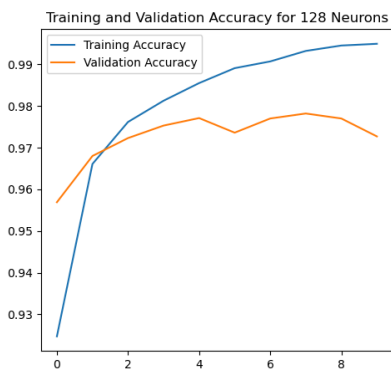
1875/1875 — 2s 1ms/step - accuracy: 0.9955 - loss: 0.0159 - val\_accuracy: 0.9727 - val\_loss: 0.0894

313/313 — 1s 953us/step - accuracy: 0.9740 - loss: 0.0882

Test Accuracy for 128-Neuron Model: 0.9782

313/313 — 0s 773us/step

	precision	recall	f1-score	support
0	0.98	0.99	0.99	980
1	0.99	0.99	0.99	1135
2	0.98	0.97	0.98	1032
3	0.97	0.99	0.98	1010
4	0.97	0.98	0.98	982
5	0.99	0.97	0.98	892
6	0.99	0.97	0.98	958
7	0.98	0.97	0.97	1028
8	0.97	0.97	0.97	974
9	0.97	0.98	0.97	1009
accuracy			0.98	10000
macro avg	0.98	0.98	0.98	10000
weighted avg	0.98	0.98	0.98	10000



# Part 4: Custom Neural Network for 99% average F1 Score

Results, graphs, reports:

Model: "sequential\_8"

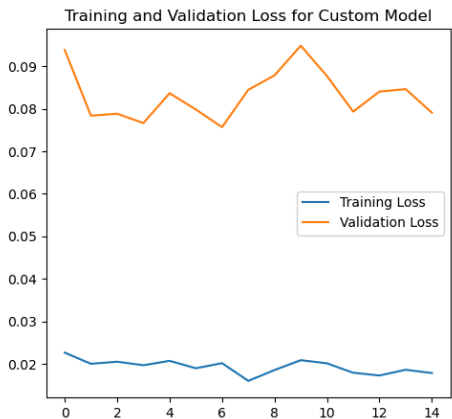
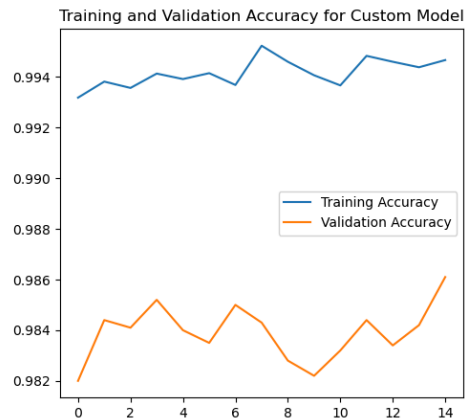
Layer (type)	Output Shape	Param #
dense_20 (Dense)	(None, 1024)	803,840
dropout_5 (Dropout)	(None, 1024)	0
dense_21 (Dense)	(None, 512)	524,800
dropout_6 (Dropout)	(None, 512)	0
dense_22 (Dense)	(None, 256)	131,328
dropout_7 (Dropout)	(None, 256)	0
dense_23 (Dense)	(None, 128)	32,896
dense_24 (Dense)	(None, 10)	1,290

Total params: 1,494,154 (5.70 MB)

Trainable params: 1,494,154 (5.70 MB)

Non-trainable params: 0 (0.00 B)

Epoch 1/15  
938/938 — 5s 4ms/step - accuracy: 0.9935 - loss: 0.0209 - val\_accuracy: 0.9820 - val\_loss: 0.0939  
Epoch 2/15  
938/938 — 3s 4ms/step - accuracy: 0.9935 - loss: 0.0216 - val\_accuracy: 0.9844 - val\_loss: 0.0784  
Epoch 3/15  
938/938 — 3s 4ms/step - accuracy: 0.9933 - loss: 0.0215 - val\_accuracy: 0.9841 - val\_loss: 0.0788  
Epoch 4/15  
938/938 — 4s 4ms/step - accuracy: 0.9940 - loss: 0.0187 - val\_accuracy: 0.9852 - val\_loss: 0.0767  
Epoch 5/15  
938/938 — 3s 4ms/step - accuracy: 0.9947 - loss: 0.0177 - val\_accuracy: 0.9840 - val\_loss: 0.0837  
Epoch 6/15  
938/938 — 3s 4ms/step - accuracy: 0.9945 - loss: 0.0182 - val\_accuracy: 0.9835 - val\_loss: 0.0798  
Epoch 7/15  
938/938 — 3s 4ms/step - accuracy: 0.9942 - loss: 0.0179 - val\_accuracy: 0.9850 - val\_loss: 0.0757  
Epoch 8/15  
938/938 — 3s 4ms/step - accuracy: 0.9957 - loss: 0.0159 - val\_accuracy: 0.9843 - val\_loss: 0.0845  
Epoch 9/15  
938/938 — 3s 4ms/step - accuracy: 0.9947 - loss: 0.0184 - val\_accuracy: 0.9828 - val\_loss: 0.0879  
Epoch 10/15  
938/938 — 4s 4ms/step - accuracy: 0.9950 - loss: 0.0170 - val\_accuracy: 0.9822 - val\_loss: 0.0948  
Epoch 11/15  
938/938 — 4s 4ms/step - accuracy: 0.9932 - loss: 0.0200 - val\_accuracy: 0.9832 - val\_loss: 0.0877  
Epoch 12/15  
938/938 — 3s 4ms/step - accuracy: 0.9947 - loss: 0.0179 - val\_accuracy: 0.9844 - val\_loss: 0.0793  
Epoch 13/15  
938/938 — 4s 4ms/step - accuracy: 0.9944 - loss: 0.0175 - val\_accuracy: 0.9834 - val\_loss: 0.0841  
Epoch 14/15  
938/938 — 4s 4ms/step - accuracy: 0.9946 - loss: 0.0185 - val\_accuracy: 0.9842 - val\_loss: 0.0846  
Epoch 15/15  
938/938 — 4s 4ms/step - accuracy: 0.9950 - loss: 0.0168 - val\_accuracy: 0.9861 - val\_loss: 0.0791

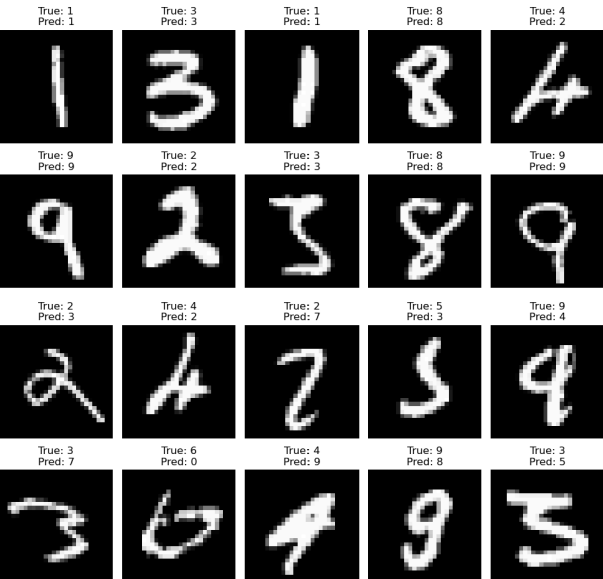
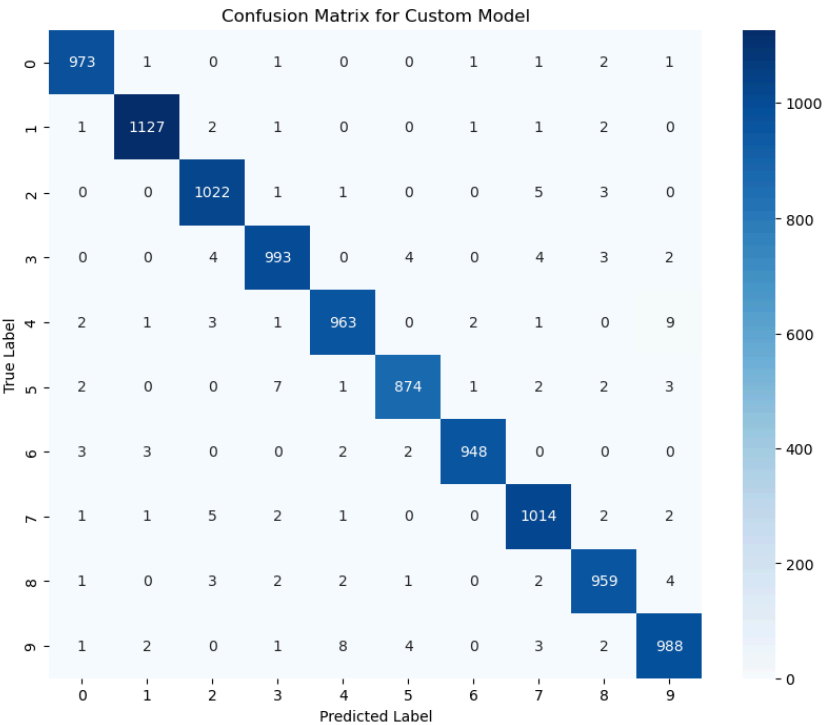


313/313 1s 1ms/step - accuracy: 0.9829 - loss: 0.1059

Test Accuracy for Custom Model: 0.9861

313/313 0s 1ms/step

	precision	recall	f1-score	support
0	0.99	0.99	0.99	980
1	0.99	0.99	0.99	1135
2	0.98	0.99	0.99	1032
3	0.98	0.98	0.98	1010
4	0.98	0.98	0.98	982
5	0.99	0.98	0.98	892
6	0.99	0.99	0.99	958
7	0.98	0.99	0.98	1028
8	0.98	0.98	0.98	974
9	0.98	0.98	0.98	1009
accuracy			0.99	10000
macro avg	0.99	0.99	0.99	10000
weighted avg	0.99	0.99	0.99	10000



## Part 5: Discussion Questions

**1. Epochs:** When neural networks are being trained, an epoch is a single run of the whole training dataset. The model uses the training data to minimize the loss function in order to update its parameters during each epoch.

Epochs are used to iteratively train the model, allowing it to gradually learn and improve its predictions across multiple cycles through the data.

Though too many epochs might cause overfitting, in which the model memorizes training data rather than generalizing, increasing the number of epochs can enhance model performance by giving it more time to learn. On the other hand, underfitting, in which the model does not learn enough from the data, might occur when there are insufficient epochs.

**2. Batch Size:** Batch size refers to the number of training samples processed before the model's internal parameters are updated. For example, in a batch size of 32, the model computes the gradient and updates the weights after processing 32 samples.

A smaller batch size allows the model to update its weights more frequently, potentially leading to faster convergence but noisier updates. Larger batch sizes provide smoother updates and are more computationally efficient but may require more memory and converge more slowly.

**3. Dropout:** Dropout is a regularization technique where a random subset of neurons is "dropped" (set to zero) during training. This prevents the model from relying too heavily on specific neurons and forces it to learn more robust features.

Dropout helps prevent overfitting by reducing the model's reliance on specific neurons, effectively encouraging it to distribute learning across the network. It also improves generalization, enabling the model to perform better on unseen data.